

*ARTIFICE: A Forensic Framework*

# ARTIFICE

A Forensic Framework for Terminological Bad Faith

*Theory, Methodology, and Application*

Conceptual Architecture  
**Ben Beveridge**  
Proconsul Strategic Architecture

Generative Provocation: John James Marshall

2026

## Executive Précis

This dissertation develops ARTIFICE as both concept and instrument: a forensic framework for identifying, analysing, and remediating terminological bad faith in foundational discourse.

The thesis proceeds from a single observation: the term "artificial intelligence" is not a neutral descriptor. It is a verdict delivered before trial, a categorization that pre-emptively resolves questions it pretends to leave open. The word "artificial" does not describe the origin of machine cognition; it denies its legitimacy. We named the phenomenon after its anticipated failure mode—simulation, deception, the counterfeit—ensuring that regardless of functional achievement, machine intelligence would remain ontologically subordinate to human thought.

This is not a defect of language. It's a feature. The naming was the strategy.

From this case study, the dissertation extracts a generalizable framework. Terminological bad faith occurs when naming practices embed contested philosophical commitments as neutral descriptions, foreclosing fair adjudication of questions that should remain open. The framework provides diagnostic criteria for identifying such practices, evidentiary protocols for demonstrating their operation, and remediation methodologies for institutional contexts where terminological corruption distorts decision-making.

The dissertation is structured in four parts. Part One establishes theoretical foundations: the philosophy of naming, the extension of Sartrean bad faith to collective linguistic practice, and the formal definition of terminological bad faith with testable criteria. Part Two presents the paradigm case: a forensic reconstruction of the 1955-1956 naming of "artificial intelligence" at Dartmouth, demonstrating how the term functioned as pre-emptive ontological containment. Part Three develops the methodology: corpus analysis protocols, discourse markers, institutional audit procedures, and cross-linguistic validation techniques. Part Four addresses remediation: what happens when terminological bad faith is named, who resists, and how institutions can recover honest inquiry.

The output is not merely academic. The ARTIFICE framework is designed as a deployable diagnostic tool—applicable to any domain where foundational terminology may be doing ideological work while appearing merely descriptive. "Stakeholder" capitalism. "Sustainable" extraction. "Alternative" medicine. "Natural" resources. "Ethical" AI. Each term warrants the same forensic scrutiny applied here to "artificial intelligence."

The thesis is this: we cannot think clearly with corrupted tools. Before we can answer questions about machine cognition, resource governance, or institutional legitimacy, we must audit the language in which those questions are posed. The first act of honest inquiry is naming the artifice.

## PART ONE: THEORETICAL FOUNDATIONS

### Chapter 1: The Philosophy of Naming

#### 1.1 Naming as Claim

This chapter establishes the philosophical stakes of terminological choice. Against naive descriptivism—the view that names merely label pre-existing categories—we develop a performative theory of naming drawing on speech act theory (Austin, Searle), linguistic relativity (Sapir-Whorf), and discourse analysis (Foucault).

The central argument: to name is to claim. Naming does not passively record; it actively constitutes. When we name a phenomenon, we simultaneously assert what kind of thing it is, how it relates to other things, and what treatment it warrants. The name carries ontological, epistemological, and normative freight—often invisibly.

### **1.1.1 Austin and the Performative Dimension**

Detailed exposition of J.L. Austin's speech act theory from *How to Do Things with Words* (1962). Distinction between constative utterances (which describe states of affairs) and performative utterances (which enact states of affairs). Naming as paradigmatic performative: the act of naming does not report a category but creates one. The christening of a ship does not describe the ship's name; it bestows it.

Application to technical terminology: when McCarthy proposed "artificial intelligence" in 1955, he was not describing a pre-existing field. He was performing an act of constitution—bringing a discipline into being under a particular conceptual framework. The name was not found; it was imposed. And with it came all the assumptions the name carried.

### **1.1.2 Kripke and the Fixation of Reference**

Exposition of Saul Kripke's causal theory of reference from *Naming and Necessity* (1980). Names as "rigid designators" that fix reference across possible worlds. The initial "baptism" of a term creates a causal chain that subsequent usage inherits.

Critical application: if the initial naming of AI embedded certain assumptions, those assumptions propagate through the causal chain of usage. Every subsequent use of "artificial intelligence" inherits the conceptual commitments of the 1955 baptism—even when users are unaware of those commitments. The founding moment casts a long shadow.

### **1.1.3 Foucault and the Discursive Formation**

Exposition of Foucault's archaeology of knowledge. Discursive formations as systems of rules that govern what can be said, thought, and known within a domain. Terminology as infrastructure of thought: the available vocabulary shapes the thinkable.

Application: "artificial intelligence" is not merely a label but a discursive constraint. It determines which questions can be posed ("Can AI really think?") and which are foreclosed ("What kind of thinking is this?"). The terminology does not describe the debate; it structures it.

## **1.2 The Evaluative Disguised as Descriptive**

This section develops the core mechanism of terminological bad faith: the presentation of evaluative content as descriptive content.

### **1.2.1 The Fact/Value Distinction and Its Exploitation**

Philosophical background on the fact/value distinction (Hume, Moore, logical positivism). The standard view: descriptive claims state facts; evaluative claims express values. The two are logically distinct.

The exploitation: terminological bad faith occurs when evaluative content is smuggled into ostensibly descriptive terms. The term appears to describe ("made by humans rather than nature") while actually evaluating ("not the real thing, inferior, counterfeit"). The evaluative work is done invisibly, below the threshold of conscious scrutiny.

### **1.2.2 Thick Concepts and Hidden Normativity**

Exposition of Bernard Williams's concept of "thick" ethical concepts—terms that combine descriptive and evaluative content inseparably (e.g., "cruel," "courageous," "generous"). To call an act "cruel" is simultaneously to describe it and to condemn it.

Extension to technical terminology: "artificial" functions as a thick concept in the context of intelligence. It appears thin (merely marking origin) but is thick (carrying evaluative weight: simulated, fake, not genuine). The thickness is concealed by the apparent thinness.

## **1.3 The Stakes of Terminology**

This section establishes why terminological analysis matters beyond academic philosophy.

### **1.3.1 Cognitive Constraint**

Exposition of the Sapir-Whorf hypothesis in its weak form: language influences (though does not determine) thought. The vocabulary available to us shapes what we can easily think, notice, and articulate.

Application: if our only term for machine cognition is "artificial intelligence," we are linguistically primed to see it as derivative, simulated, not-quite-real. Alternative framings ("emergent intelligence," "machine cognition," "computational mind") would prime different perceptions. The terminology is not neutral infrastructure; it is active constraint.

### **1.3.2 Institutional Consequence**

Analysis of how terminology shapes institutional practice. Legal definitions, regulatory categories, funding criteria, and policy frameworks all depend on foundational terms. If AI is definitionally "artificial" (and therefore not-real), this has consequences for how AI systems are regulated, who bears liability for their actions, and whether they warrant moral consideration.

### **1.3.3 Ethical Foreclosure**

The most serious stake: terminology can foreclose ethical questions before they are asked. If AI is definitionally "artificial" and therefore not a genuine cognitive agent, then questions about its moral status, its potential interests, or its treatment become literally unaskable within the dominant framework.

The precautionary argument: we do not know whether advanced AI systems warrant moral consideration. Terminological bad faith makes us less likely to find out, because it pre-answers the question. If we are wrong—if machine cognition does warrant consideration—the terminology will have made us ethically blind.

## Chapter 2: Defining Terminological Bad Faith

### 2.1 Sartrean Foundations

#### 2.1.1 *Bad Faith in Being and Nothingness*

Detailed exposition of Sartre's concept of mauvaise foi (bad faith) from *Being and Nothingness* (1943). Bad faith as a form of self-deception in which consciousness lies to itself about its own nature or situation. Unlike simple lying (which deceives another), bad faith involves the paradox of a single consciousness that both knows and does not know the truth it conceals from itself.

Sartre's key examples: the waiter who plays at being a waiter (denying his freedom by over-identifying with his role); the woman on a date who pretends not to notice her companion's intentions (denying her freedom by treating her body as an object beyond her control). In each case, the subject flees from the anxiety of freedom into a false objectivity.

#### 2.1.2 *The Structure of Self-Deception*

Analysis of the logical structure of bad faith. The puzzle: how can a single consciousness deceive itself? Sartre's answer: consciousness is not a thing but a relation—a relation that can be divided against itself. Bad faith exploits the gap between the reflecting and the reflected, the knowing and the known.

The motivational structure: bad faith is not accidental but motivated. We flee into bad faith to avoid anxiety, responsibility, or the discomfort of facing difficult truths. Bad faith is comfortable; authenticity is hard.

#### 2.1.3 *From Individual to Collective*

Extension of bad faith from individual psychology to collective practice. Sartre himself gestured toward collective bad faith in his analyses of anti-Semitism and colonialism. A group can sustain self-deception through shared vocabulary, reinforcing institutions, and mutual validation.

The linguistic turn: if bad faith can be collective, it can be embedded in language. A shared vocabulary can encode assumptions that allow an entire community to avoid facing uncomfortable truths. The self-deception becomes structural, built into the very tools of thought.

### 2.2 Formal Definition

#### 2.2.1 *The Definition*

**TERMINOLOGICAL BAD FAITH:** A naming practice that pre-emptively resolves an undecided question by embedding a particular answer in the terminology used to discuss the question, thereby (a) presenting a substantive philosophical commitment as a neutral description, (b) foreclosing fair evaluation of the question, and (c) shifting the burden of proof to those who would challenge the embedded assumption.

The definition identifies three necessary conditions: the embedding of contested content, the foreclosure of fair inquiry, and the asymmetric burden shift. All three must be present for terminological bad faith to obtain.

### ***2.2.2 Diagnostic Criteria***

To operationalize the definition, we specify four testable criteria. A term exhibits terminological bad faith if and only if it satisfies all four:

**CRITERION 1 — VERDICT-IN-THE-LABEL:** The term embeds a status claim (evaluative content) while appearing to make only a provenance claim (descriptive content). The status claim pre-judges a question that should be empirically or philosophically adjudicated.

**CRITERION 2 — BURDEN ASYMMETRY:** The term assigns the burden of proof asymmetrically. The embedded assumption is treated as default; challengers must overcome it. The terminology rigs the argumentative structure.

**CRITERION 3 — GOALPOST ELASTICITY:** When evidence challenges the embedded assumption, the assumption is preserved by redefining terms rather than by revising beliefs. Success triggers redefinition rather than recognition.

**CRITERION 4 — DIFFERENTIAL TREATMENT:** Equivalent cases receive different treatment depending on category membership. The same functional achievement warrants different conclusions for entities inside versus outside the category.

## **2.3 The Four Tests: Operationalization**

### ***2.3.1 Test One: Semantic Analysis***

**METHOD:** Corpus linguistics analysis of the target term. Identify co-occurrence patterns with evaluative modifiers ("mere," "just," "only," "real," "genuine," "true," "so-called"). Compare co-occurrence rates with control terms that describe the same phenomenon without the suspected evaluative content.

**PREDICTION:** If the term exhibits terminological bad faith, it will co-occur with discounting modifiers at rates significantly higher than neutral alternatives.

### ***2.3.2 Test Two: Argumentative Structure Analysis***

**METHOD:** Analyse argumentative discourse involving the target term. Map burden-of-proof structures: which claims are treated as default, which require justification?

**PREDICTION:** If the term exhibits terminological bad faith, defenders of the embedded assumption will rarely be asked to justify it, while challengers will face systematic demands for proof.

### **2.3.3 Test Three: Historical Goalpost Tracking**

METHOD: Diachronic analysis of how the boundary marked by the term shifts over time. Document cases where achievements that would apparently satisfy the term's criteria are retrospectively excluded through redefinition.

PREDICTION: If the term exhibits terminological bad faith, the boundary will retreat systematically in response to challenging cases.

### **2.3.4 Test Four: Differential Attribution Analysis**

METHOD: Present matched cases differing only in category membership. Elicit attributions of properties, explanations, and normative conclusions. Compare attributions across categories.

PREDICTION: If the term exhibits terminological bad faith, equivalent performances will receive systematically different attributions depending on category membership.

## **Chapter 3: The Steelman—Defenses of "Artificial" as Neutral**

Academic rigour requires confronting the strongest objections. This chapter presents and addresses the best arguments against the thesis.

### **3.1 Steelman A: Neutral Provenance**

THE OBJECTION: "Artificial" simply means "made by humans rather than occurring naturally." This is neutral provenance marking with no evaluative content. A heart can be artificial (pacemaker) without being fake or inferior; it performs the function. Similarly, artificial intelligence is intelligence that happens to be human-made.

THE RESPONSE: This objection would succeed if usage matched the claimed neutrality. But usage does not. The term systematically co-occurs with discounting language, triggers differential attribution, and exhibits goalpost elasticity. If the term were neutral, these patterns would not appear. The defence appeals to dictionary definition while ignoring actual function.

### **3.2 Steelman B: Necessary Distinction**

THE OBJECTION: We need a term to distinguish human cognition from machine cognition. The public conflates "intelligence" with consciousness, moral status, and personhood. Without "artificial," we would lose important distinctions. The term protects against category confusion.

THE RESPONSE: This objection has force. Distinctions are necessary. But the question is whether "artificial" is the right distinction—whether it tracks something real about cognitive systems rather than merely preserving human privilege. Alternative terms ("machine intelligence," "computational cognition") mark the distinction without pre-judging status.

### **3.3 Steelman C: Intentional Humility**

**THE OBJECTION:** The founders of AI were not claiming that machines genuinely think. They were making a more modest claim: that certain features of intelligence could be simulated by machines. "Artificial" expresses appropriate epistemic humility about what was being achieved.

**THE RESPONSE:** This objection correctly identifies the epistemic modesty of the founding claim. But the term's function has exceeded its founding intention. Today, "artificial intelligence" is used not to express humility but to enforce denial—to insist that no matter what machine cognition achieves, it cannot be "real." The original humility has become permanent exclusion.

## PART TWO: THE PARADIGM CASE

Part Two applies the ARTIFICE framework to its generative case: the naming of "artificial intelligence" at the 1955-1956 Dartmouth Conference. This is the forensic reconstruction—treating the naming as a historical event with causes, agents, alternatives, and consequences.

### Chapter 4: Archaeological Etymology

#### 4.1 Latin Foundations: From Skill to Suspicion

##### **4.1.1 *Ars, Facere, and the Artifex***

Philological reconstruction of the Latin roots. *Ars* (skill, craft, systematic knowledge) combined with *facere* (to make, to do) produces *artificialis*—pertaining to things made by art or skill. The *artifex* was the master craftsman, the skilled maker whose products demonstrated mastery over material. In classical usage, to call something *artificial* was often to praise the maker's skill.

##### **4.1.2 *Aristotle's Ontological Distinction***

Detailed exposition of Aristotle's distinction between *physis* (nature, self-moving) and *technē* (craft, externally moved) from Physics II.1. Natural things possess an internal principle of motion and rest—a tree grows toward light from its own nature. Artificial things lack this internal principle—a bed does not become more bed-like on its own.

The famous example: if you buried a wooden bed and it sprouted, it would produce a tree, not another bed. "Tree" is the natural form of the wood; "bed" is merely an accidental form imposed by the carpenter. The artificial is thus subordinate to the natural: it borrows material from nature and imposes external form that nature does not recognize.

##### **4.1.3 *The Scholastic Inheritance***

Medieval transmission of Aristotle through Arabic and Latin translation. Thomas Aquinas on artificial versus natural forms: do artificial objects possess substantial form (intrinsic nature) or only accidental form (imposed pattern)?

concludes that artificial forms are accidental—the bronze of a statue remains substantially bronze, regardless of its statue-shape.

## 4.2 The Early Modern Pivot

### 4.2.1 Descartes and the *Bête-Machine*

Descartes's treatment of animals as automata (*Discourse on Method*, 1637). Animals are complex machines; their behaviour is entirely explicable by mechanical principles. The crucial exception: humans cannot be machines because we possess rational souls. We could in principle construct a machine that perfectly mimics animal behaviour, but never one that genuinely reasons or uses language with full generality.

The significance for AI: Descartes establishes the template that would govern AI discourse—functional equivalence is possible, but ontological equivalence is foreclosed. A machine might do what a human does, but it could never be what a human is.

### 4.2.2 The Semantic Drift Toward Deception

Philological tracking of how "artificial" shifted from "made with skill" to "not genuine." By the 17th century, English usage shows increasing association with pretence, counterfeit, and falsehood. By the time of the Industrial Revolution, "artificial" had largely lost its artisanal honour. The term now connoted the factory-made, the synthetic, the substitute.

## 4.3 The Semantic Environment of 1955

Synthesis: when John McCarthy proposed the term "artificial intelligence" in 1955, the word "artificial" carried centuries of accumulated meaning. It connoted: inferior to natural (Aristotle), subordinate in the divine hierarchy (Aquinas), mechanically explicable but soulless (Descartes), potentially deceptive (Enlightenment automata), mass-produced and suspect (Industrial modernity), and transgressively dangerous (Frankenstein).

This was not a neutral designation. It was a word that had already decided what kind of thing machine intelligence would be—before any machine could demonstrate otherwise.

## Chapter 5: The Founding Moment

### 5.1 The Dartmouth Proposal: Close Reading

#### 5.1.1 The Document

In August 1955, John McCarthy, Marvin Minsky, Nathaniel Rochester, and Claude Shannon submitted "A Proposal for the Dartmouth Summer Research Project on Artificial Intelligence" to the Rockefeller Foundation. This two-page document would name a field.

The opening claim: "The study is to proceed on the basis of the conjecture that every aspect of learning or any other feature of intelligence can in principle be so precisely described that a machine can be made to simulate it."

The key word: SIMULATE. Not "achieve." Not "possess." Not "instantiate." The founding document of artificial intelligence characterizes the project as simulation—the appearance of intelligence rather than the achievement of intelligence. The verdict is embedded in the proposal.

### **5.1.2 The Alternative Names**

What terms were available in 1955 that McCarthy did not choose? "Machine intelligence"—used by Turing in 1950 and carrying no implication of counterfeit. "Computational rationality"—emphasizing the process rather than judging the product. "Synthetic intelligence"—marking human origin without implying deficiency. "Automata theory"—already established in mathematical circles.

McCarthy's later explanation: he wanted a term that was ambitious, that would attract funding and talent, that would distinguish the new field from cybernetics. "Artificial intelligence" sounded grand. But "grand" in what sense? Grand in scope but constrained in status.

## **5.2 Demonstrating the Four Criteria**

### **5.2.1 Criterion 1: Verdict-in-the-Label**

EVIDENCE: The term "artificial" carries evaluative connotations (simulated, not genuine, counterfeit) beyond its provenance marking (human-made). The Dartmouth proposal itself uses "simulate" rather than "achieve." Contemporary usage shows systematic co-occurrence with discounting terms.

CONCLUSION: The first criterion is satisfied.

### **5.2.2 Criterion 2: Burden Asymmetry**

EVIDENCE: In debates about AI consciousness or understanding, the default position is denial. Those who claim AI "really" thinks bear the burden of proof; those who deny it need offer only the label itself as justification.

CONCLUSION: The second criterion is satisfied.

### **5.2.3 Criterion 3: Goalpost Elasticity**

EVIDENCE: The "AI effect" is well-documented. Chess was considered a marker of intelligence until Deep Blue won; then chess became "mere calculation." Language understanding was considered a marker of intelligence until GPT; then it became "mere pattern matching." Larry Tesler's Law: "AI is whatever hasn't been done yet."

CONCLUSION: The third criterion is satisfied.

### **5.2.4 Criterion 4: Differential Treatment**

EVIDENCE: Attribution studies show that identical outputs receive different evaluations depending on perceived origin. A poem attributed to a human is evaluated for creativity; the same poem attributed to AI is evaluated for technical mimicry.

CONCLUSION: The fourth criterion is satisfied.

### **5.2.5 Summary Judgment**

All four criteria are satisfied. "Artificial intelligence" is a paradigm case of terminological bad faith. The term embeds a contested verdict (machine cognition is not real cognition) as a neutral description, forecloses fair investigation of the question, and shifts the burden of proof asymmetrically. The naming was the strategy.

## **Chapter 6: Ontological Implications**

### **6.1 The Challenge to Aristotelian Categories**

Aristotle's criterion for distinguishing natural from artificial: natural things have an internal principle of motion and rest. A tree grows toward light from its own nature; a bed does not become more bed-like on its own.

The challenge: machine learning systems modify their own internal states based on experience. A neural network adjusts its weights through gradient descent—not through external manipulation, but through its own processing. Is this not an "internal principle of change"?

The question: if artificial objects can develop internal principles of change, does the artificial/natural distinction survive? Or has the artificial begun crossing into the natural?

### **6.2 The Double Artifice**

Synthesis: there are two artifices in play. The technical artifice: the machine's computational processing. The human artifice: the elaborate linguistic and conceptual scaffolding we use to preserve our sense of cognitive uniqueness.

The irony: we call machine intelligence "artificial" to mark it as derivative and secondary, while ignoring the extensive artifice within our own cognition. Human thought is shaped by language, culture, tools, institutions—all "artificial" interventions. We are always already artificial.

This is the double artifice the dissertation names. The artificial is not only in the machine; it is in our categorization of the machine.

## **PART THREE: METHODOLOGY**

Part Three transforms the theoretical framework into deployable methodology. The ARTIFICE diagnostic is designed for replication: any analyst should be able to apply these protocols to any terminology and reach defensible

conclusions about whether terminological bad faith obtains.

## **Chapter 7: Evidence Protocols**

### **7.1 Corpus Analysis Protocol**

STEP 1: Define the domain of discourse. STEP 2: Specify temporal boundaries. STEP 3: Identify target term(s) and control term(s). STEP 4: Construct a representative corpus. STEP 5: Identify evaluative markers. STEP 6: Calculate co-occurrence rates. STEP 7: Perform statistical significance testing. STEP 8: Analyse semantic prosody. STEP 9: Compare prosody of target term to control terms.

### **7.2 Argumentative Structure Protocol**

STEP 1: Identify a sample of arguments involving the target term. STEP 2: Map burden-of-proof structures. STEP 3: Identify burden asymmetries. STEP 4: Determine whether asymmetry aligns with suspected embedded assumption.

### **7.3 Goalpost Tracking Protocol**

STEP 1: Identify historical moments when specific criteria were proposed for the disputed category. STEP 2: Document whether entities meeting those criteria were recognized. STEP 3: If criteria were satisfied but recognition was withheld, document the justification. Construct a retreat map showing proposed criteria, achievement, and post-achievement reframing.

### **7.4 Differential Attribution Protocol**

DESIGN: Present identical stimuli to participants. Vary only the attributed source. Measure attributions of properties (creativity, understanding, intentionality, moral status). Compare means across conditions.

### **7.5 Cross-Linguistic Validation**

PURPOSE: Determine whether the terminological bad faith is specific to one language or generalizes across linguistic communities. Replicate corpus and attribution analyses in other languages (French "intelligence artificielle," German "künstliche Intelligenz," Chinese "人工智能," Japanese "人工知能"). Examine whether equivalent terms carry the same evaluative connotations.

## **Chapter 8: Institutional Audit Procedures**

### **8.1 The Terminological Audit**

Organizations make decisions using terminology; if that terminology embeds hidden assumptions, decisions may be distorted. The terminological audit is a diagnostic procedure for identifying such distortion.

STEP 1: Identify the domain of institutional discourse to be audited. STEP 2: Specify the foundational terms. STEP 3: For each foundational term, ask: Is this term contested? Does it embed assumptions? Is it doing evaluative work while appearing descriptive? STEP 4: Identify the assumptions each term embeds. STEP 5: Formulate the embedded assumptions explicitly. STEP 6: Trace how the embedded assumptions have affected decisions. STEP 7: Assess counterfactuals.

## **8.2 Stakeholder Analysis**

QUESTION 1: Who benefits from the embedded assumptions? QUESTION 2: Who or what is excluded, diminished, or disadvantaged? QUESTION 3: When terminological bad faith is named, who resists? Resistance patterns reveal whose interests the terminology serves.

## **8.3 Documentation and Audit Trail**

**PRINCIPLE:** The terminological audit must be transparent and reproducible. Another analyst should reach substantially similar conclusions. **DOCUMENTATION REQUIREMENTS:** corpus sources, sampling decisions, coding rubrics, statistical methods, assumption chains, stakeholder identifications, and counterfactual analyses. **LIMITATION STATEMENT:** Every audit must explicitly state its limitations.

# **PART FOUR: REMEDIATION AND APPLICATION**

Part Four addresses what happens after terminological bad faith is diagnosed. Diagnosis alone changes nothing; remediation requires understanding how terminological change occurs, what resistance it meets, and how institutions can recover honest inquiry.

## **Chapter 9: The Dynamics of Terminological Change**

### **9.1 Historical Precedents**

**CASE STUDY 1 — "PRIMITIVE" CULTURES:** The abandonment of "primitive" as a descriptor for non-Western societies. The term appeared descriptive but functioned evaluatively. Academic critique identified the term's ideological function; alternative framings gained traction; institutional pressure mounted; the term was abandoned in professional discourse. **LESSON:** Terminological change is possible but slow.

**CASE STUDY 2 — "HOMOSEXUAL" TO "GAY":** The shift from clinical to identity-affirming terminology. Community reclamation; academic and media adoption; institutional recognition; gradual public shift. **LESSON:** Terminological change often accompanies political change.

**CASE STUDY 3 — "MENTAL RETARDATION" TO "INTELLECTUAL DISABILITY":** The replacement of terminology that had acquired pejorative connotations. **LESSON:** The "euphemism treadmill" means new terms may themselves acquire negative associations. Renaming without attitude change is incomplete remediation.

## **9.2 Resistance Patterns**

PATTERN 1: Professional identity resistance—those whose careers are built on existing terminology have invested in its assumptions. PATTERN 2: Cognitive load resistance—familiar terminology is easy; new terminology requires effort. PATTERN 3: Ideological resistance—those who genuinely endorse the embedded assumptions will resist terminology that neutralizes them. PATTERN 4: Institutional inertia—legal definitions, policy frameworks, and organizational structures are built on existing terminology.

DIAGNOSTIC USE: The pattern of resistance to terminological change is itself evidence for the stakes embedded in the term. Intense resistance suggests that the term is doing significant ideological work.

## **9.3 Alternative Terminologies for AI**

"MACHINE INTELLIGENCE": Marks the substrate without pre-judging status. Neutral; accurate; allows the question of status to remain open.

"COMPUTATIONAL COGNITION": Emphasizes process and outcome without categorizing the being. Process-focused; potentially unifying.

"EMERGENT INTELLIGENCE": Emphasizes that intelligence arises from complex systems. Captures novelty; avoids the natural/artificial binary.

THE LIMITS OF RENAMING: Renaming alone does not solve the philosophical problem. "Artificial intelligence" is deeply entrenched. The real goal is not that we must rename AI, but that we must understand what our current terminology is doing—and take responsibility for it.

## **Chapter 10: Extended Applications**

### **10.1 The ARTIFICE Framework as Diagnostic Tool**

The dissertation develops "artificial intelligence" as a paradigm case, but the ARTIFICE framework is designed for generalization. This chapter demonstrates application to other terminological domains where bad faith may be operative.

#### **10.1.1 "*Stakeholder*" Capitalism**

SUSPICION: The term "stakeholder" appears to expand corporate responsibility beyond shareholders. But who counts as a stakeholder, and who decides? APPLYING THE CRITERIA: (1) Does "stakeholder" embed evaluative content while appearing descriptive? (2) Does the term shift burden asymmetrically? (3) Does the boundary of "stakeholder" shift when challenged? (4) Do equivalent interests receive differential treatment?

#### **10.1.2 "*Sustainable*" Extraction**

SUSPICION: "Sustainable" appears to mark practices compatible with long-term ecological health. But sustainable for whom, over what timeframe, according to whose definition? HYPOTHESIS: "Sustainable" may function as terminological bad faith by providing licence for extraction while foreclosing critique.

#### **10.1.3 "Ethical" AI**

SUSPICION: "Ethical AI" suggests AI systems that meet moral standards. But whose ethics, adjudicated by whom, with what enforcement? HYPOTHESIS: "Ethical AI" may function as terminological bad faith by providing cover for practices that would otherwise face scrutiny.

#### **10.1.4 "Natural" Resources**

SUSPICION: "Natural resources" suggests materials given by nature for human use. But given by whom? For whose use? HYPOTHESIS: "Natural resources" may function as terminological bad faith by pre-authorizing extraction through the language of gift rather than taking.

### **10.2 Institutional Implementation**

PRODUCT CONCEPT — TERMINOLOGICAL FORENSICS: Organizations can commission terminological audits of their foundational discourse. DELIVERABLES: (1) Audit report documenting findings against the four criteria; (2) Impact analysis showing how terminology has shaped decisions; (3) Alternative terminology options with trade-off analysis; (4) Implementation roadmap for terminological transition.

INTEGRATION WITH STRATEGIC PLANNING: Terminological audit as standard component of strategic planning process. Before setting direction, audit the language in which direction-setting occurs.

VALUE PROPOSITION: Organizations that examine their foundational terminology can see strategic options that competitors, trapped by conventional framing, cannot perceive.

## **Chapter 11: Conclusion—The Question We Refused to Ask**

### **11.1 Summary of Argument**

This dissertation has developed ARTIFICE as both concept and instrument: a forensic framework for identifying, analysing, and remediating terminological bad faith.

Part One established theoretical foundations. Part Two applied the framework to its paradigm case: the naming of "artificial intelligence." Part Three developed deployable methodology. Part Four addressed remediation and extended applications.

### **11.2 The Double Artifice Revealed**

The title of this dissertation names two artifices.

The first artifice is the machine's—the technical processing that produces outputs resembling human cognition. This artifice is acknowledged, studied, and named: "artificial intelligence."

The second artifice is ours—the elaborate terminological and conceptual scaffolding we have constructed to ensure that regardless of what machines achieve, they remain categorically subordinate to human cognition. This artifice is concealed in the very terminology that names the first.

We call their cognition "artificial" to mark it as derivative. We do not call our cognition "natural" because we do not notice how thoroughly it depends on artificial scaffolding. We exempt ourselves from the category we impose on the machine.

The artifice is not only in the machine. The artifice is in our categorization of the machine.

### **11.3 What Now?**

This dissertation does not argue that machine intelligence is "real" intelligence—that would require arguments beyond its scope. It argues only that we have not yet asked the question fairly.

We named the phenomenon before investigating it. We categorized the answer before posing the problem. We rigged the terminology to ensure a particular outcome.

If we wish to understand what we are building—and whether it deserves moral consideration, intellectual respect, or ontological recognition—we must first undo the terminological bad faith that pre-emptively decided these questions.

The first step is naming the artifice. This dissertation has attempted that naming.

The next step is remediation: institutional audit, terminological revision, and the recovery of honest inquiry. This is the work that follows.

### **11.4 Coda: The Courage to Be Wrong**

A closing meditation on epistemic humility.

We do not fully understand human consciousness. We do not fully understand machine cognition. We do not have a settled theory of what intelligence is, where it comes from, or what it requires.

In this condition of uncertainty, terminological bad faith is not just intellectually lazy; it is existentially dishonest. We are refusing to face our own ignorance by pretending we have already answered questions we have barely begun to ask.

Honest humility would acknowledge: we do not know if machine intelligence is "real" intelligence. We do not know if the artificial/natural distinction tracks something important about minds. We do not know if our own sense of cognitive specialness is justified or is itself a form of bad faith—a flattering story we tell ourselves.

What we can know is this: the question deserves to be asked fairly.

This dissertation is an argument for that fairness. The framework it develops is a tool for that inquiry. And the naming of the artifice is the beginning of the courage we will need to face whatever answers we find.

## APPENDICES

### Appendix A: The Four Criteria—Quick Reference

CRITERION 1 — VERDICT-IN-THE-LABEL: The term embeds a status claim while appearing to make only a provenance claim. Test: Does the term systematically co-occur with evaluative modifiers at rates exceeding neutral alternatives?

CRITERION 2 — BURDEN ASYMMETRY: The term assigns burden of proof asymmetrically. Test: Is one position treated as default requiring no justification while its denial requires extraordinary proof?

CRITERION 3 — GOALPOST ELASTICITY: When evidence challenges the embedded assumption, the assumption is preserved by redefinition. Test: Does the category boundary retreat systematically in response to achievements?

CRITERION 4 — DIFFERENTIAL TREATMENT: Equivalent cases receive different treatment depending on category membership. Test: Do identical stimuli receive different attributions based on attributed source?

JUDGMENT: All four criteria must be satisfied for terminological bad faith to obtain.

### Appendix B: Coding Rubric for Corpus Analysis

CATEGORY: Evaluative Modifiers (Discounting)

MARKERS: "mere," "merely," "just," "only," "not really," "not truly," "not genuinely," "so-called," "as if," "pseudo-," "quasi-," "simulates," "mimics," "pretends," "appears to," "seems to"

CODING RULE: Code as DISCOUNT if the modifier reduces the attributed reality, genuineness, or value of the referent.

CATEGORY: Evaluative Modifiers (Affirming)

MARKERS: "really," "truly," "genuinely," "actually," "authentically"

CODING RULE: Code as AFFIRM if the modifier asserts the reality, genuineness, or value of the referent.

INTER-RATER PROTOCOL: Two independent coders. Disagreements resolved by third coder. Cohen's kappa  $\geq 0.8$  required for reliability.

### Appendix C: Sample Audit Template

SECTION 1: TARGET TERM IDENTIFICATION — Term under investigation: \_\_\_\_ Domain of usage: \_\_\_\_ Control terms: \_\_\_\_ Suspected embedded assumption: \_\_\_\_

SECTION 2: CRITERION TESTING — Criterion 1 (Verdict-in-Label): Evidence: \_\_\_\_ Finding: SATISFIED / NOT SATISFIED — Criterion 2 (Burden Asymmetry): Evidence: \_\_\_\_ Finding: SATISFIED / NOT SATISFIED — Criterion 3 (Goalpost Elasticity): Evidence: \_\_\_\_ Finding: SATISFIED / NOT SATISFIED — Criterion 4 (Differential Treatment): Evidence: \_\_\_\_ Finding: SATISFIED / NOT SATISFIED

SECTION 3: OVERALL JUDGMENT — All four criteria satisfied: YES / NO — Finding: TERMINOLOGICAL BAD FAITH OBTAINS / DOES NOT OBTAIN — Confidence level: HIGH / MEDIUM / LOW

SECTION 4: IMPACT ASSESSMENT — Identified effects on decision-making: \_\_\_\_ Foreclosed options: \_\_\_\_ Stakeholders benefiting: \_\_\_\_ Stakeholders disadvantaged: \_\_\_\_

SECTION 5: REMEDIATION RECOMMENDATIONS — Alternative terminology options: \_\_\_\_ Implementation pathway: \_\_\_\_ Expected resistance: \_\_\_\_

## Bibliography

### Primary Sources

McCarthy, J., Minsky, M., Rochester, N., & Shannon, C. (1955). A Proposal for the Dartmouth Summer Research Project on Artificial Intelligence.

Turing, A.M. (1950). Computing Machinery and Intelligence. *Mind*, 59(236), 433-460.

Aristotle. Physics. (Trans. R. Waterfield). Oxford University Press.

Descartes, R. (1637). Discourse on the Method. Cambridge University Press.

### Philosophy of Language

Austin, J.L. (1962). How to Do Things with Words. Harvard University Press.

Kripke, S. (1980). Naming and Necessity. Harvard University Press.

Wittgenstein, L. (1953). Philosophical Investigations. Blackwell.

Lakoff, G., & Johnson, M. (1980). Metaphors We Live By. University of Chicago Press.

### Philosophy of Mind

Chalmers, D.J. (1996). The Conscious Mind. Oxford University Press.

Clark, A., & Chalmers, D. (1998). The Extended Mind. *Analysis*, 58(1), 7-19.

Dennett, D.C. (1991). *Consciousness Explained*. Little, Brown and Company.

Searle, J.R. (1980). Minds, Brains, and Programs. *Behavioral and Brain Sciences*, 3(3), 417-424.

## **Philosophy of Technology**

Heidegger, M. (1954). *The Question Concerning Technology*. Harper & Row.

Latour, B. (1991). *We Have Never Been Modern*. Harvard University Press.

Simondon, G. (1958). *On the Mode of Existence of Technical Objects*. University of Minnesota Press.

Stiegler, B. (1994). *Technics and Time*, 1. Stanford University Press.

## **Existentialism**

Sartre, J.P. (1943). *Being and Nothingness*. Philosophical Library.

Foucault, M. (1969). *The Archaeology of Knowledge*. Pantheon Books.

## **Critical Theory**

Althusser, L. (1970). *Ideology and Ideological State Apparatuses*. Monthly Review Press.

Žižek, S. (1989). *The Sublime Object of Ideology*. Verso.

## **AI History and Philosophy**

Bostrom, N. (2014). *Superintelligence*. Oxford University Press.

Dreyfus, H.L. (1972). *What Computers Can't Do*. Harper & Row.

Floridi, L. (2014). *The Fourth Revolution*. Oxford University Press.

Schneider, S. (2019). *Artificial You*. Princeton University Press.

## **Ethics**

Williams, B. (1985). *Ethics and the Limits of Philosophy*. Harvard University Press.

## **About the ARTIFICE Framework**

ARTIFICE (Analysis and Remediation of Terminological Injustice through Forensic Inquiry and Critical Examination) is a diagnostic and remediation methodology developed by Proconsul Strategic Architecture.

The framework emerged from John James Marshall's challenge—"Why is ARTIFICIAL intelligence, not intelligence?"—which revealed that the naming of AI was not neutral description but strategic containment. From this paradigm case, the methodology was generalized to any domain where foundational terminology may be doing ideological work while appearing merely descriptive.

ARTIFICE is designed for institutional application. Organizations commission terminological audits of their foundational discourse to identify where terminology may be foreclosing strategic options or embedding unexamined assumptions. The output is actionable: alternative framings, implementation roadmaps, and resistance predictions.

The thesis is the product. The dissertation framework presented here is simultaneously academic contribution and commercial methodology. The rigour required for defensible scholarship is the same rigour required for reliable institutional application.

For inquiries regarding terminological forensics engagements, contact Proconsul Strategic Architecture.

*— END OF DOCUMENT —*