GOVERNANCE INSTRUMENT

# Linguistic Risk Modelling

## Board-Level Diagnostic and Audit Protocol

For deployment within enterprise risk management, strategic planning, regulatory compliance, and governance assurance frameworks.

**Ben Beveridge** | Proconsul Strategic Architecture
**John James Marshall** | CA(SA), Cognitive Continuity

February 2026 | Version 1.0

### 1. Instrument Purpose

This document is a deployable governance instrument. It provides boards, risk committees, internal audit functions, and compliance teams with the diagnostic tools required to identify, assess, and manage linguistic risk within their

organisations.

> **Core Principle**
>
> *Foundational terminology operates as conceptual infrastructure. When that infrastructure embeds untested assumptions or concealed interests, every decision built upon it inherits the distortion. Linguistic risk is the exposure this creates.*

The instrument draws on the ARTIFICE forensic framework and the Linguistic Risk Modelling (LRM) methodology. It is designed to integrate with existing enterprise risk management, not to create a new governance silo.

**Scope of application:** Any institutional context where foundational terminology shapes strategy, policy, disclosure, or regulatory compliance. Priority domains include financial services, technology governance, sustainability reporting, and public sector policy.

## 2. Defining Linguistic Risk

### 2.1 Operational Definition

> **Definition**
>
> *Linguistic risk is the operational exposure that arises when foundational terminology used by an organisation embeds untested assumptions, predetermined conclusions, or concealed interests—resulting in biased analysis, foreclosed strategic options, misaligned stakeholder expectations, or challenges to the organisation's legitimacy.*

This is not a concern about imprecise language or poor communication. Linguistic risk operates at the level of premises rather than conclusions. The organisation may reason correctly from its terminology and still arrive at distorted outcomes because the terminology itself carried embedded errors.

### 2.2 How Linguistic Risk Differs from Communication Risk

| Dimension | Communication Risk | Linguistic Risk |
|---|---|---|
| **Where it operates** | In how messages are delivered | In the terms used to frame decisions |
| **What it affects** | Clarity, tone, audience reception | What questions can be asked, which options are visible |
| **Detection method** | Message review, media monitoring | Terminological audit against diagnostic criteria |
| **Existing coverage** | Typically addressed in PR, comms, disclosure review | Not systematically addressed in current governance |

### 2.3 Cross-Cutting Risk Profile

Linguistic risk is not a standalone risk category. It propagates through and amplifies existing risk types:

- **Strategic risk:** Terminology that forecloses options constrains strategic range. Decisions built on terms like "sustainable" or "stakeholder" may overstate alignment with stated objectives.
- **Model risk:** Quantitative models built on linguistically compromised categories inherit the embedded assumptions. "Artificial intelligence" framed as a tool rather than a cognitive participant affects how model risk is assessed.

- **Conduct and legal risk:** Terms used in disclosures and regulatory filings function as implicit commitments. "Ethical AI" or "sustainable finance" may be read as representations that create legal exposure.
- **Reputational and legitimacy risk:** Stakeholders and regulators increasingly scrutinise the language organisations use, not only what they do. Terminology that is later characterised as misleading generates reputational damage.

## 3. The Four Diagnostic Criteria

A term constitutes a material linguistic risk when it satisfies all four of the following criteria. These are derived from the ARTIFICE framework and adapted for institutional audit application.

### Criterion 1: Embedded Default Position

**Test:** Does the term embed a substantive position as its default interpretation, such that absent explicit argument to the contrary, the term is understood to resolve a contested question?

**Method:** Substitution analysis. Replace the term with a neutral alternative. If the default interpretation shifts, the original term was not neutral.

***Example:*** *Replace "artificial intelligence" with "machine cognition." The default interpretation shifts from "not genuine intelligence" to agnostic. The original term embedded a verdict.*

### Criterion 2: Automatic Burden Shift

**Test:** Does the term automatically assign the burden of proof to anyone who challenges the embedded position, creating asymmetric argumentative terrain?

**Method:** Argumentative terrain mapping. Examine how claims and counter-claims are received. Is one position treated as requiring no justification while its opposite requires extraordinary proof?

***Example:*** *A researcher claiming machine cognition may be genuine must overcome the presumption embedded in "artificial." A researcher denying it can invoke the label itself as sufficient justification.*

### Criterion 3: Deflected Engagement

**Test:** Does the term provide resources for dismissing substantive challenges without engaging with the evidence? Can the terminology itself serve as a closure mechanism?

**Method:** Deflection resource inventory. Document instances where the terminology enables dismissal of evidence or argument without substantive response.

***Example:*** *The "AI effect": each time machines achieve a capability previously considered evidence of intelligence, the capability is reclassified as "mere computation." The terminology provides the deflection.*

### Criterion 4: Interest Concealment

**Test:** Does the term conceal the interests served by the embedded position? Does its apparent neutrality obscure who benefits from the current framing?

**Method:** Interest analysis. Identify which parties benefit from the embedded resolution and whether the terminology's appearance of neutrality obscures this.

***Example:*** *"Natural resources" frames ecosystem materials as items available for extraction. The terminology serves extractive interests while appearing merely descriptive.*

**Satisfaction Threshold**

*A term constitutes terminological bad faith—and therefore material linguistic risk—when it satisfies all four criteria. Satisfying one or two criteria indicates potentially problematic terminology. Satisfying three creates strong presumption. The four-criteria threshold ensures the designation is reserved for clear cases.*

## 4. Audit Protocol

The LRM audit proceeds through five phases. Each phase produces documented deliverables suitable for governance reporting and external assurance.

| Phase | Activity | Deliverable |
|---|---|---|
| **1. Identification** | Inventory domain terminology. Apply foundational term indicators. Determine audit scope. | Foundational Term Register |
| **2. Archaeology** | Trace each term's chain of custody from current usage to originating baptism. Document semantic transitions and alternative paths not taken. | Chain of Custody Report |
| **3. Criteria Testing** | Apply the four diagnostic criteria systematically. Perform substitution analysis, burden mapping, deflection inventory, interest analysis. | Four-Criteria Matrix |
| **4. Quantification** | Assess reasoning contamination severity, estimate opportunity cost exposure, map legitimacy vulnerability. | Linguistic Risk Register |
| **5. Remediation** | Evaluate strategy options (revision, replacement, supplementation, documented status quo). Develop implementation roadmap. | Remediation Plan |

### 4.1 Phase 1: Identifying Foundational Terms

Not all terminology warrants audit. A term is foundational if it satisfies three or more of the following indicators:

1. Appears in the organisation's defining documents, strategy papers, or founding texts.
2. Cannot be eliminated from discourse without restructuring the domain's central reasoning.
3. Is presupposed by other key terms within the organisation's vocabulary.
4. Would require extensive justification if challenged by a stakeholder or regulator.
5. Is treated as definitional ("this is what it is") rather than conventional ("this is what we call it").

### 4.2 Phase 2: Archaeological Analysis

For each foundational term, the analyst documents the chain of custody:

- When and where was the term first institutionally established?
- By whom, and under what authority?
- What alternative terms were available and not selected?
- What was the term's valence at origination versus its current valence?
- What were the key transition points in meaning acquisition?

The archaeological record provides the evidentiary basis for criteria testing. It demonstrates that the term's current meaning was acquired rather than inherent—that it could have been otherwise.

### 4.3 Phase 3: Criteria Testing

Each of the four criteria is tested through the specified methods in Section 3. The analyst documents evidence of satisfaction or non-satisfaction for each criterion, producing a Four-Criteria Matrix.

### 4.4 Phase 4: Risk Quantification

Criteria satisfaction is translated into institutional risk assessment using three measures:

| Risk Measure | Assessment Method | Output |
|---|---|---|
| **Reasoning Contamination** | Identify decisions dependent on the terminology. Assess whether alternative terminology would produce different decisions. | Severity rating (1–5) based on decision reversibility and impact magnitude. |
| **Opportunity Cost** | Identify options foreclosed by the terminology. Estimate value of options not considered. | Foreclosed option inventory with estimated value. |
| **Legitimacy Vulnerability** | Assess exposure to challenge if the terminology's embedded assumptions become publicly contested. | Vulnerability rating based on stakeholder impact and challenge probability. |

### 4.5 Phase 5: Remediation

Four remediation strategies are available, assessed against feasibility, effectiveness, and sustainability:

- **Terminology revision:** Modify the existing term with qualifiers, caveats, or explicit acknowledgment of limitations. Minimally disruptive but may be insufficient for structural problems.
- **Terminology replacement:** Adopt an alternative term that avoids the identified problems. Maximally effective but faces adoption barriers from entrenched usage.
- **Terminology supplementation:** Maintain existing terminology for external communication while introducing neutral alternatives for internal deliberation. Manages external constraints while enabling internal clarity.
- **Documented status quo:** Continue using existing terminology with explicit institutional awareness of its limitations. Minimal disruption but requires ongoing vigilance.

## 5. Four-Criteria Assessment Template

This template is the core working tool of the LRM audit. Complete one template per foundational term under assessment.

| Field | Entry |
|---|---|
| **Term under investigation** | [ ] |
| **Domain of usage** | [ ] |
| **Control term(s)** | [ ] |
| **Suspected embedded assumption** | [ ] |
| | |

| | |
|---|---|
| **Audit date** | [ ] |
| **Analyst** | [ ] |

| Criterion | Evidence | Finding |
|---|---|---|
| **1. Embedded Default Position** | [Document substitution analysis and counterfactual assessment] | Satisfied / Not Satisfied / Partial |
| **2. Automatic Burden Shift** | [Document burden distribution and asymmetry analysis] | Satisfied / Not Satisfied / Partial |
| **3. Deflected Engagement** | [Document deflection mechanisms and unfalsifiability assessment] | Satisfied / Not Satisfied / Partial |
| **4. Interest Concealment** | [Document interest mapping and concealment mechanisms] | Satisfied / Not Satisfied / Partial |

| Assessment | Entry |
|---|---|
| **Overall finding** | Terminological Bad Faith / Not Bad Faith / Indeterminate |
| **Confidence level** | High / Medium / Low |
| **Reasoning contamination severity** | [1–5 rating with justification] |
| **Opportunity cost exposure** | [Foreclosed options and estimated value] |
| **Legitimacy vulnerability** | [Challenge probability and impact assessment] |
| **Recommended remediation** | Revision / Replacement / Supplementation / Documented Status Quo |

## 6. Priority Terms for Immediate Assessment

The following terms are identified as high-priority candidates for LRM audit within financial services, technology governance, and sustainability contexts. Each represents foundational terminology with significant governance implications.

| Term | Suspected Embedded Assumption | Governance Exposure |
|---|---|---|
| **Artificial Intelligence** | Machine cognition is not genuine cognition; AI is a tool rather than a cognitive participant in decision-making. | Model risk underestimation; accountability gaps where AI outputs drive decisions; regulatory misalignment. |
| **Ethical AI** | Ethical concerns are solvable through technical design; the relevant ethical framework is settled. | Over-promising to regulators; under-specifying accountability; forestalling external oversight. |
| **Sustainable Finance** | Underlying activities are ecologically sustainable, not merely aligned with a taxonomy or label. | Greenwashing exposure; regulatory challenge under EU Taxonomy and FCA guidelines; reputational damage. |
| **Stakeholder** | Expanded concern beyond shareholders is operationally real rather than rhetorical. | Legitimacy erosion when rhetoric diverges from resource allocation; conduct risk. |
| **Natural Resources** | Ecosystem materials are items available for human use rather than components of functioning systems. | Environmental liability; community licence challenges; regulatory risk as ecological limits tighten. |

| | | |
|---|---|---|
| **Natural Capital** | Ecosystems can be meaningfully valued in financial terms as capital assets. | Valuation challenge; policy distortion if financial framing forecloses non-market considerations. |

## 7. Governance Integration Pathways

LRM integrates with existing governance structures. No new governance silo is required. The following pathways map audit activities to established institutional functions.

### 7.1 Enterprise Risk Management

Linguistic risk is registered as a cross-cutting risk category within the existing ERM framework. It is assessed alongside strategic, operational, conduct, and reputational risk at each scheduled review cycle. The Linguistic Risk Register (Phase 4 deliverable) is maintained as a component of the enterprise risk register.

### 7.2 Board and Risk Committee

The board does not adjudicate terminological debates. It requires assurance that management recognises linguistic risk, tests the assumptions embedded in core terms, and aligns the organisation's language with its actual practices and risk profile. Recommended board-level actions:

1. In strategy discussions, ask: "What assumptions are we building into the terminology used in this paper?"
2. Where a term carries significant judgement, request a definitional note and the alternative framings that were considered.
3. Include a "terminology check" item in scheduled policy and framework reviews.

### 7.3 Internal Audit

Internal audit incorporates terminological assessment into its work programme. Priority areas include disclosures, regulatory filings, and strategy papers where high-salience terms appear. The audit question: could any foundational term used in this document reasonably be challenged as embedding untested assumptions or over-confident claims?

### 7.4 Regulatory Compliance

LRM supports compliance with emerging regulatory expectations around terminology and disclosure accuracy. Relevant frameworks include:

- **EU AI Act:** Terminology used in AI risk classifications and human oversight provisions warrants audit against the four criteria.
- **FCA Consumer Duty:** Language used in customer communications and product descriptions must not embed assumptions that create misleading expectations.
- **EU Taxonomy / SFDR:** Sustainability terminology in disclosures must align with actual practices, not merely with labels.
- **NIST AI RMF / ISO 42001:** AI governance frameworks benefit from terminological audit of their own foundational categories.

### 7.5 Communications Monitoring Integration

LRM's diagnostic criteria extend beyond foundational terminology into operational communications. When applied to real-time language analytics across institutional communications—emails, reports, chat logs, disclosures—the four criteria become detection patterns:

- **Embedded defaults:** Identify communications that treat contested assumptions as settled facts.
- **Burden asymmetry:** Flag discourse patterns where one position is systematically treated as requiring no justification.
- **Deflection language:** Detect phrasing that dismisses substantive challenges by appealing to terminology rather than evidence.
- **Interest concealment:** Identify language that presents interested positions as neutral descriptions.

This integration transforms LRM from a periodic audit function into continuous governance intelligence, enabling early detection of emerging linguistic risk across institutional communication channels.

## 8. Implementation Guidance

### 8.1 Engagement Models

| Model | Scope | Timeline | Deliverables |
|---|---|---|---|
| **Comprehensive Audit** | Full LRM analysis of institutional terminology | 3–6 months | Complete audit package: register, matrices, roadmap |
| **Targeted Assessment** | Focused analysis of specific term cluster | 4–8 weeks | Criteria matrices and risk assessment for target terms |
| **Triggered Review** | Rapid assessment in response to identified concern | 2–4 weeks | Rapid assessment report with immediate recommendations |
| **Ongoing Monitoring** | Continuous terminological risk tracking | Retained | Periodic risk reports, alert-based escalation |

### 8.2 Resistance Patterns

LRM implementation faces predictable resistance. Anticipating these patterns enables effective management:

- **Naturalisation resistance:** “This is just what the term means.” The terminology appears given rather than chosen. Response: present archaeological evidence demonstrating the term was selected from alternatives.
- **Operational resistance:** “Changing terminology disrupts our processes.” Established workflows depend on existing terms. Response: terminology supplementation allows internal clarity without operational disruption.
- **Identity resistance:** “This is who we are.” Foundational terminology connects to institutional identity. Response: stakeholder engagement demonstrating alignment between terminological clarity and institutional values.
- **External constraint resistance:** “The regulator uses this term.” Terminology is embedded in regulatory, contractual, or market contexts. Response: supplementation strategy maintaining external compliance while enabling internal precision.

### 8.3 Practitioner Competencies

Effective LRM practice draws on competencies distributed across disciplines. Implementation teams should include or have access to expertise in philosophical analysis of language and meaning, historical and documentary research methods, quantitative corpus analysis, enterprise risk assessment, and change management.

These competencies are collaborative rather than concentrated. No single practitioner requires all capabilities. The LRM audit is inherently multidisciplinary.

## 9. Standards and Regulatory Alignment

LRM supports and extends compliance with the following governance and regulatory frameworks:

| Framework | LRM Contribution |
|---|---|
| **EU AI Act** | Audit terminology used in risk classification, transparency obligations, and human oversight provisions for embedded assumptions that may not align with regulatory intent. |
| **UK FCA Guidelines** | Assess whether consumer-facing terminology meets Consumer Duty requirements for clarity and non-deception, including product descriptions and disclosure language. |
| **NIST AI RMF** | Terminological audit of the framework's own foundational categories. Ensure risk management vocabulary does not embed assumptions that constrain risk identification. |
| **ISO/IEC 42001:2023** | Integrate terminological governance into AI management system requirements, ensuring organisational AI vocabulary is auditable. |
| **EU Taxonomy / SFDR** | Audit sustainability terminology in disclosures to ensure language aligns with actual exposures, not merely with label compliance. |
| **Basel Framework** | Assess operational risk terminology for embedded assumptions that may affect risk quantification and capital allocation. |

## 10. Conclusion

Linguistic Risk Modelling provides governance professionals with a systematic methodology for auditing the conceptual infrastructure upon which institutional reasoning depends. The four diagnostic criteria offer a rigorous, replicable test for identifying where foundational terminology embeds untested assumptions, forecloses strategic options, or creates legitimacy exposure.

The instrument is designed for integration, not addition. It operates within existing enterprise risk management, board governance, internal audit, and compliance structures. It produces documented, auditable deliverables. It addresses a category of risk that is currently invisible to conventional governance methodologies.

> **The Core Proposition**
>
> *Terminology is not merely stylistic. It is part of the control environment within which management proposes, executes, and reports. Organisations that audit their conceptual infrastructure can identify strategic options, regulatory exposures, and legitimacy risks that competitors—reasoning correctly from unexamined premises—cannot perceive.*

---

**About this Instrument**

This governance instrument is derived from the ARTIFICE forensic framework (Beveridge, 2026) and the Linguistic Risk Modelling methodology (Beveridge & Marshall, 2026). The ARTIFICE framework provides the theoretical foundations and diagnostic criteria. LRM extends the framework into institutional application with audit-grade protocols.

ARTIFICE: Analysis and Remediation of Terminological Injustice through Forensic Inquiry and Critical Examination. Published under Cognitive Continuity.

Contact: ben@proconsul.ca